

NanoJPEG CUDA

Parallelization of a JPEG Decompression Library on a CUDA GPGPU

Lorenzo CHIOLA, s287911



**Politecnico
di Torino**

JPEG/JFIF Format

- JPEG: Image Compression

- choice of **color spaces**

- choice of **compression quality**

- choice of **subsampling**

- JFIF: Container Format

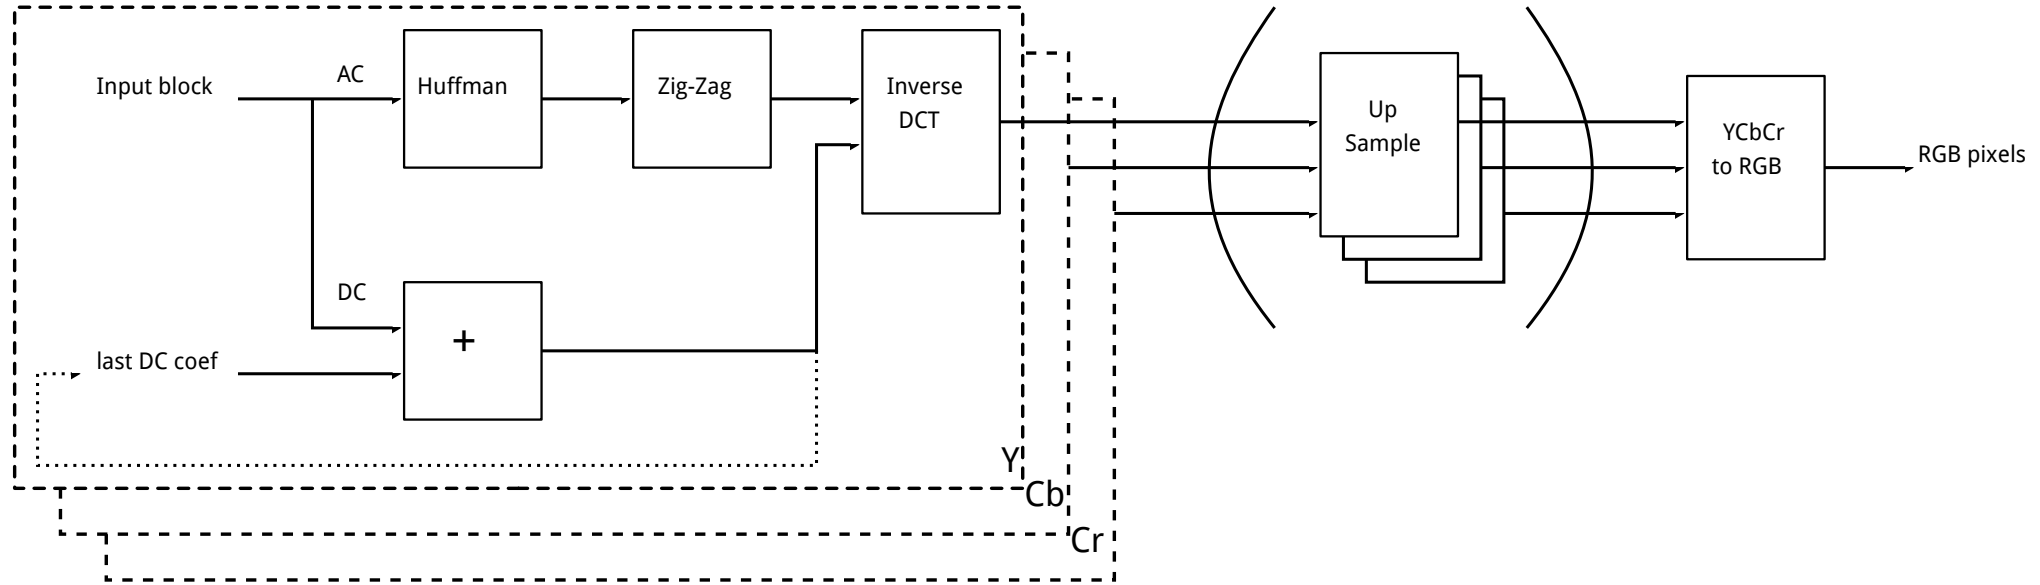
- choice of **sections** (some are mandatory)

- progressive JPEG**

JPEG Compression Overview

- 8x8 px blocks
- Chroma Subsampling
- DCT: Dual Cosine Transform
- Huffman Encoding
- Separate DC component

JPEG Compression Flow



JFIF (JPEG File Interchange Format)

- **SOF** (Start Of Frame)

Width, height, color space...

- **DHT** (Define Huffman Table)

Huffman code tree. Y DC, Y AC, Cb/Cr DC, Cb/Cr AC

- **DQT** (Define Quantization Table)

Separate for Y and Cb, Cr

- **DRI** (Define Restart Interval)

DC DPCM restart component

- (Start of) **Scan**

Actual data: 8x8 blocks

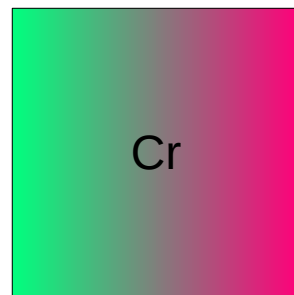
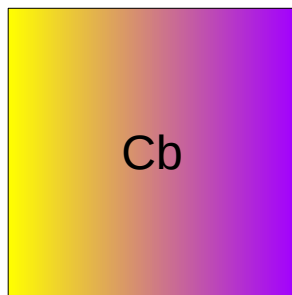
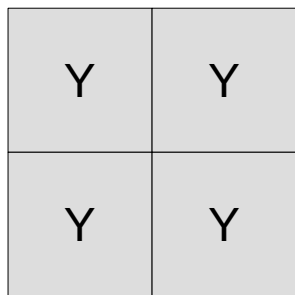
JFIF Scan

- **8x8 Blocks** grouped into MCUs (Minimum Coded Units)

RGB: {R, G, B}, {R, G, B}, ...
YcbCr 4:4:4: {Y, Cb, Cr}, {Y, Cb, Cr}, ...

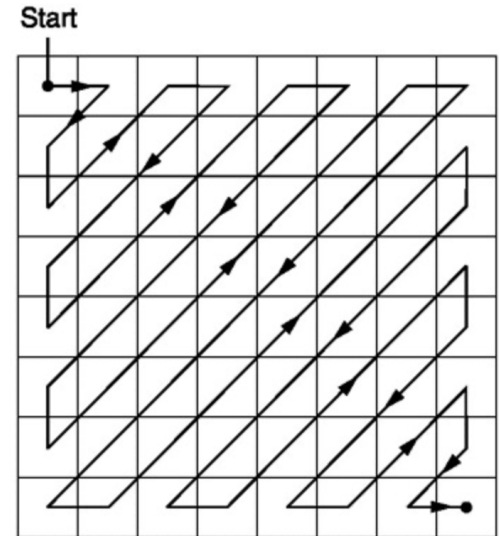
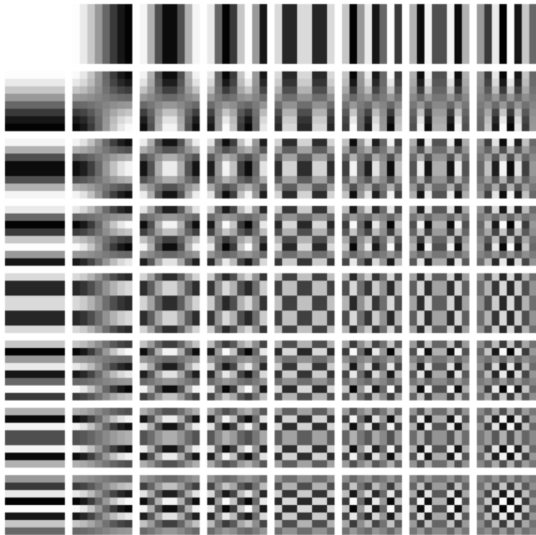
- **Chroma Subsampling** (less chrominance information)

YcbCr 4:2:2: {Y, Y, Cb, Cr}, {Y, Y, Cb, Cr}, ...
YcbCr 4:2:0: {Y, Y, Y, Y, Cb, Cr}, {Y, Y, Y, Y, Cb, Cr}, ...

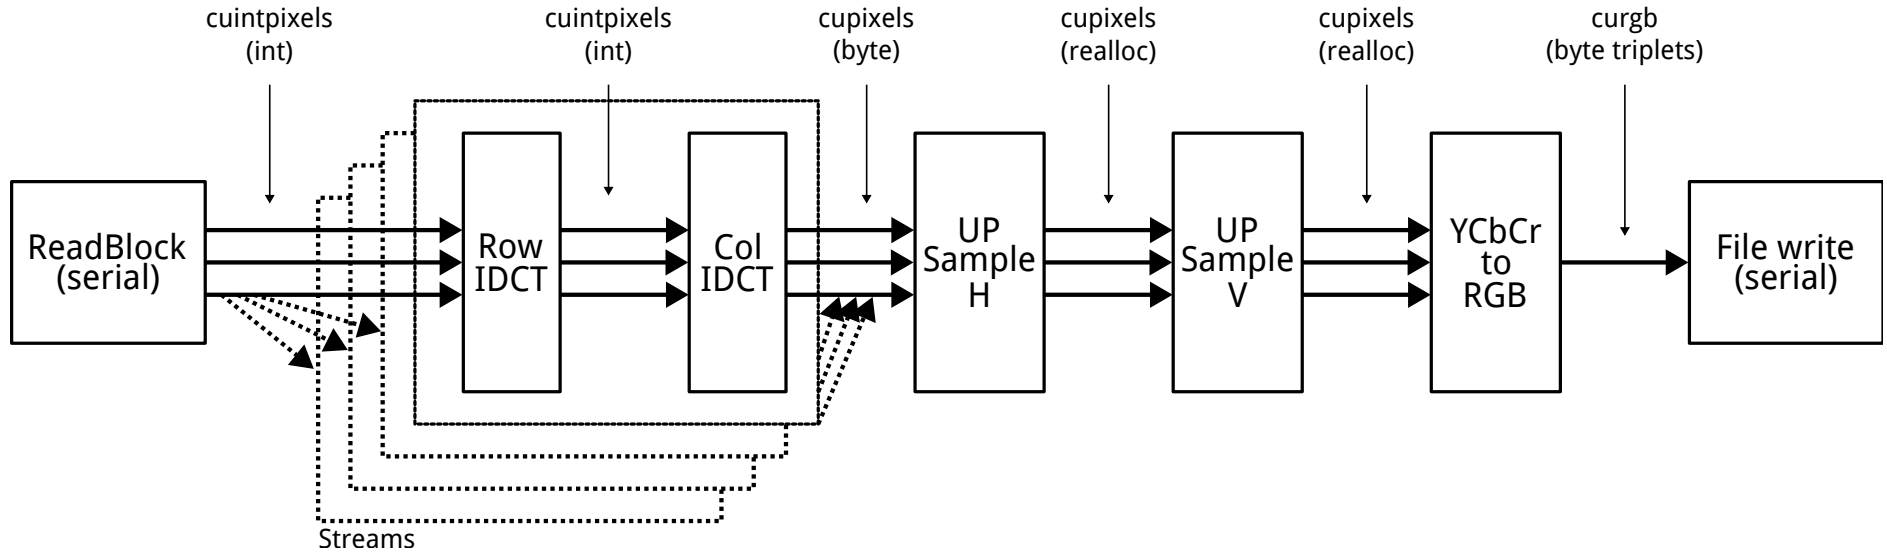


JFIF Scan: Blocks

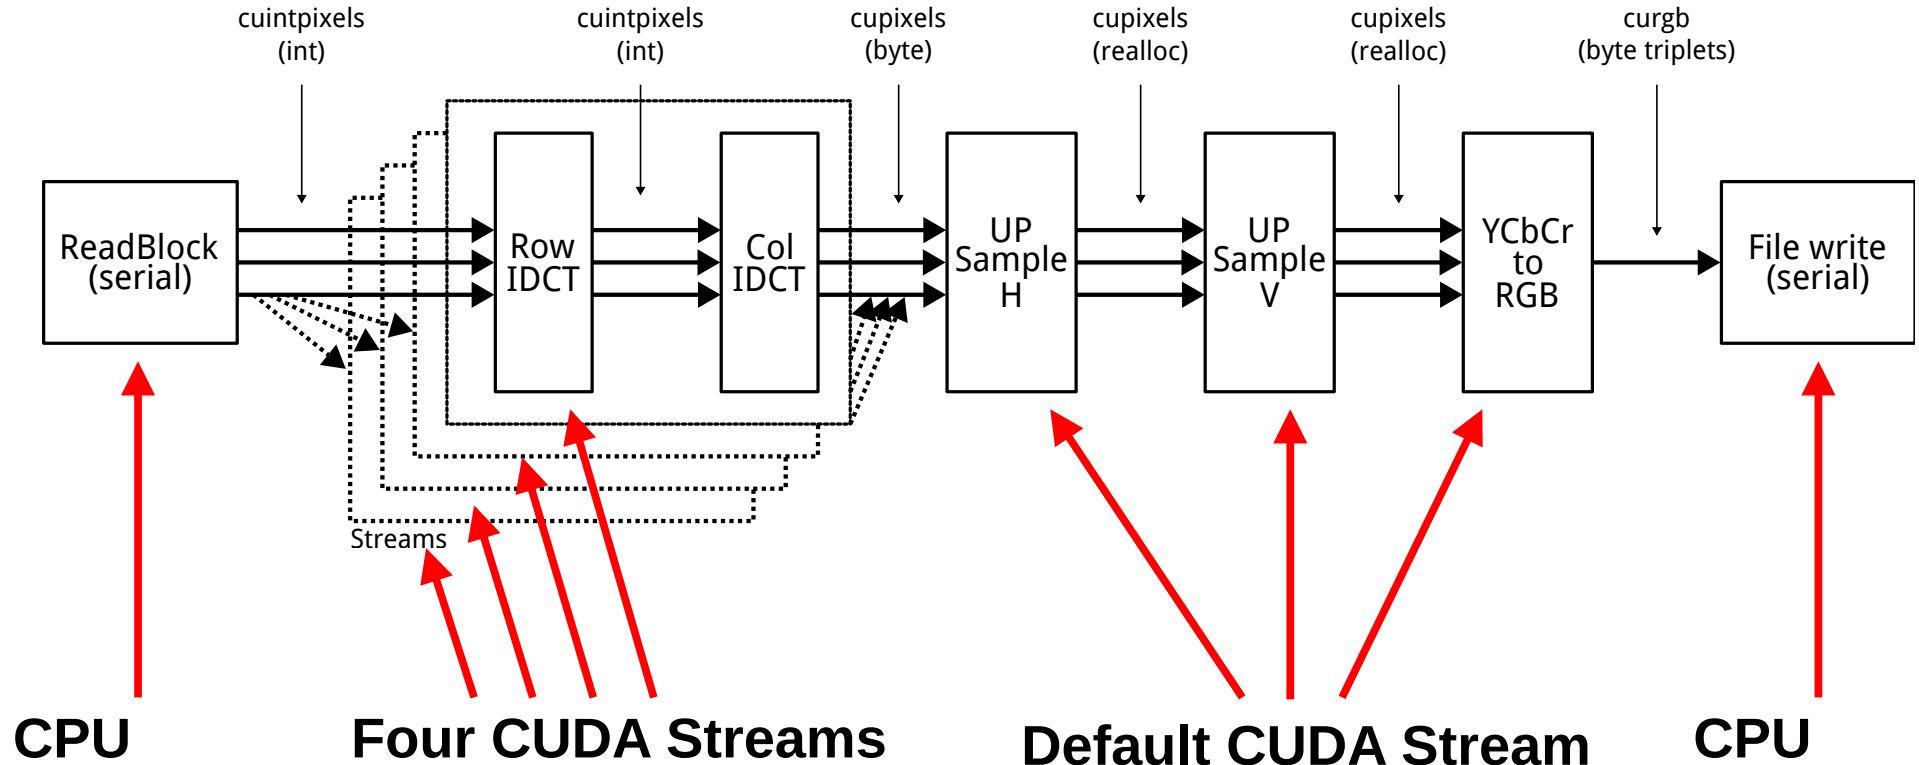
- **DC coefficient** (zero frequency: cell [0, 0])
Differential Pulse Code Modulation (DPCM)
- **AC coefficient** (the other 63 cells)
Zig-Zag order: from low to high frequency
Huffman encoded



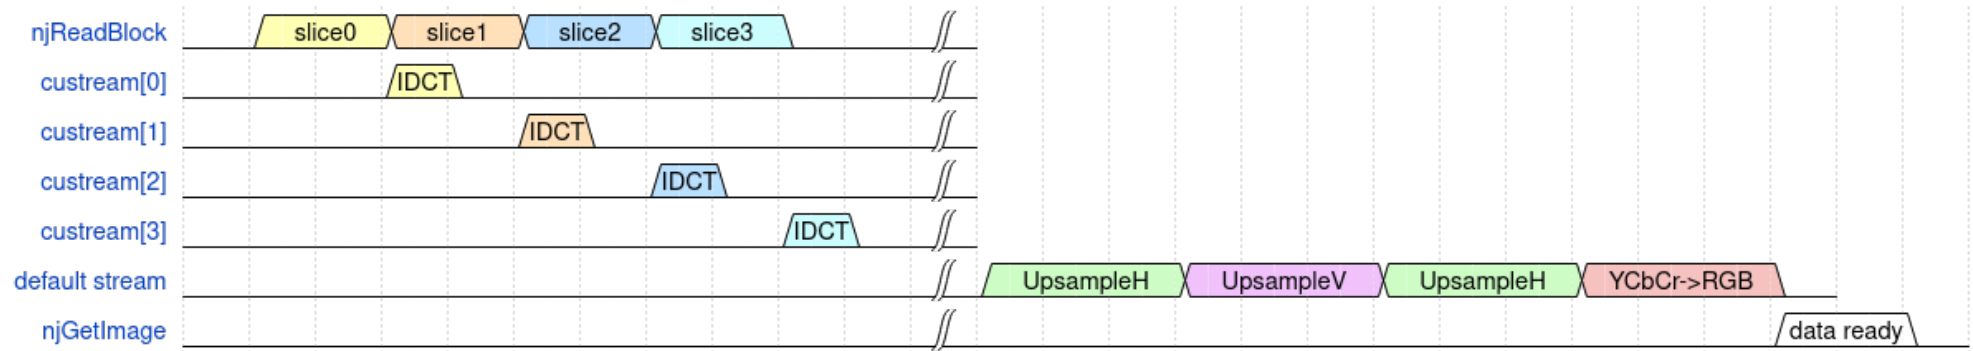
NanoJPEG CUDA Data Flow



Streamed Execution

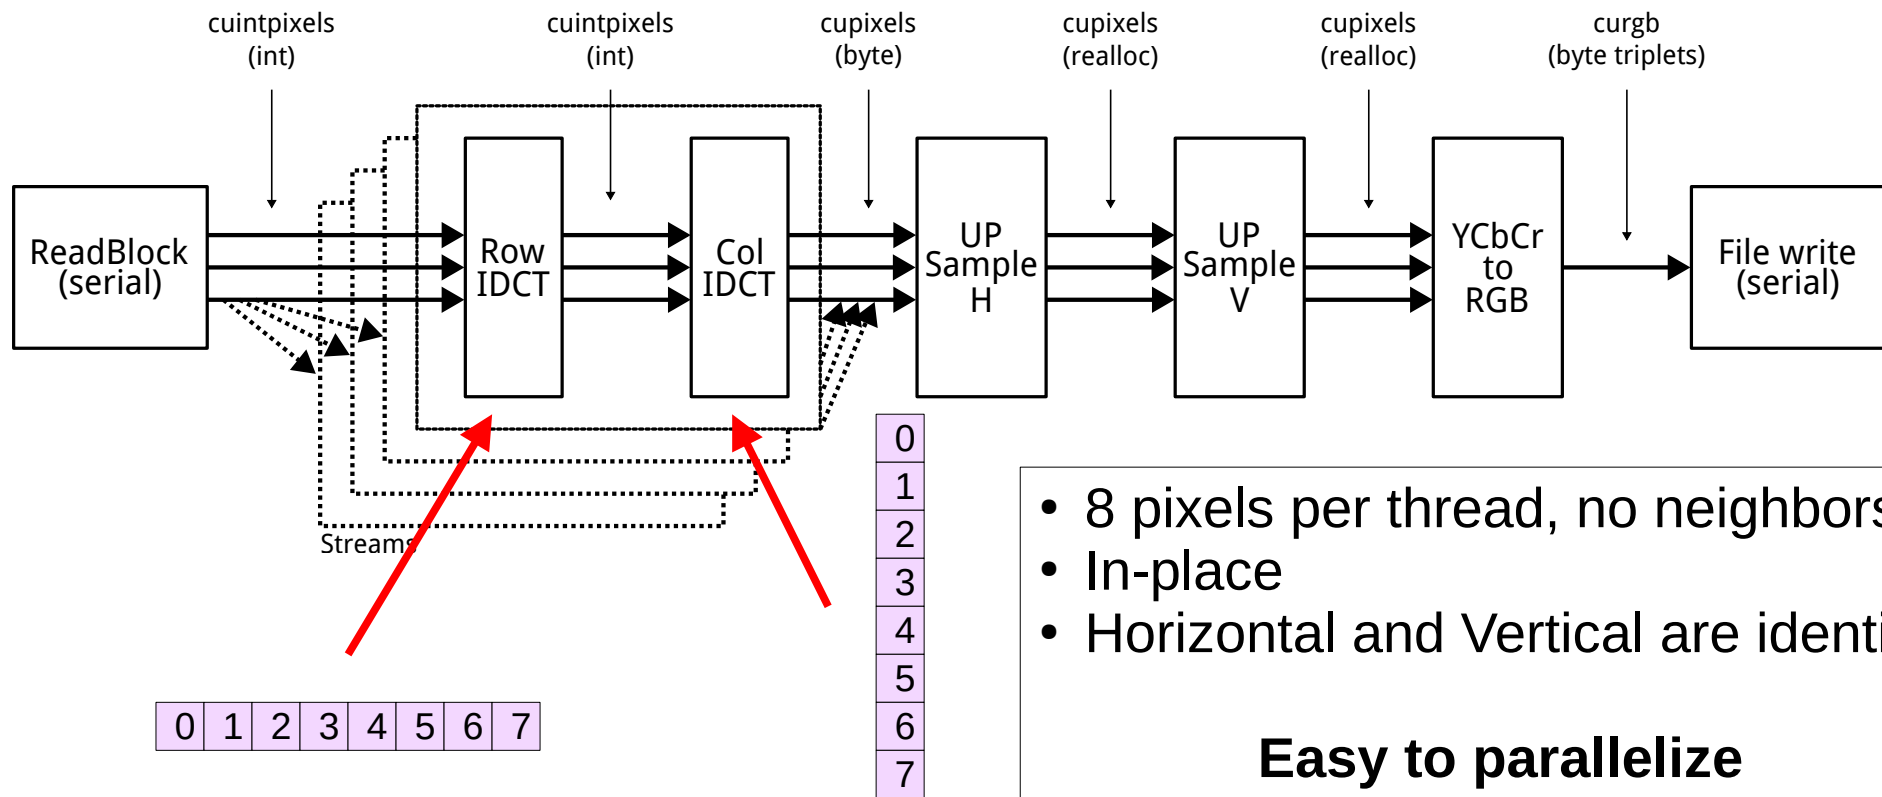


Streamed Execution

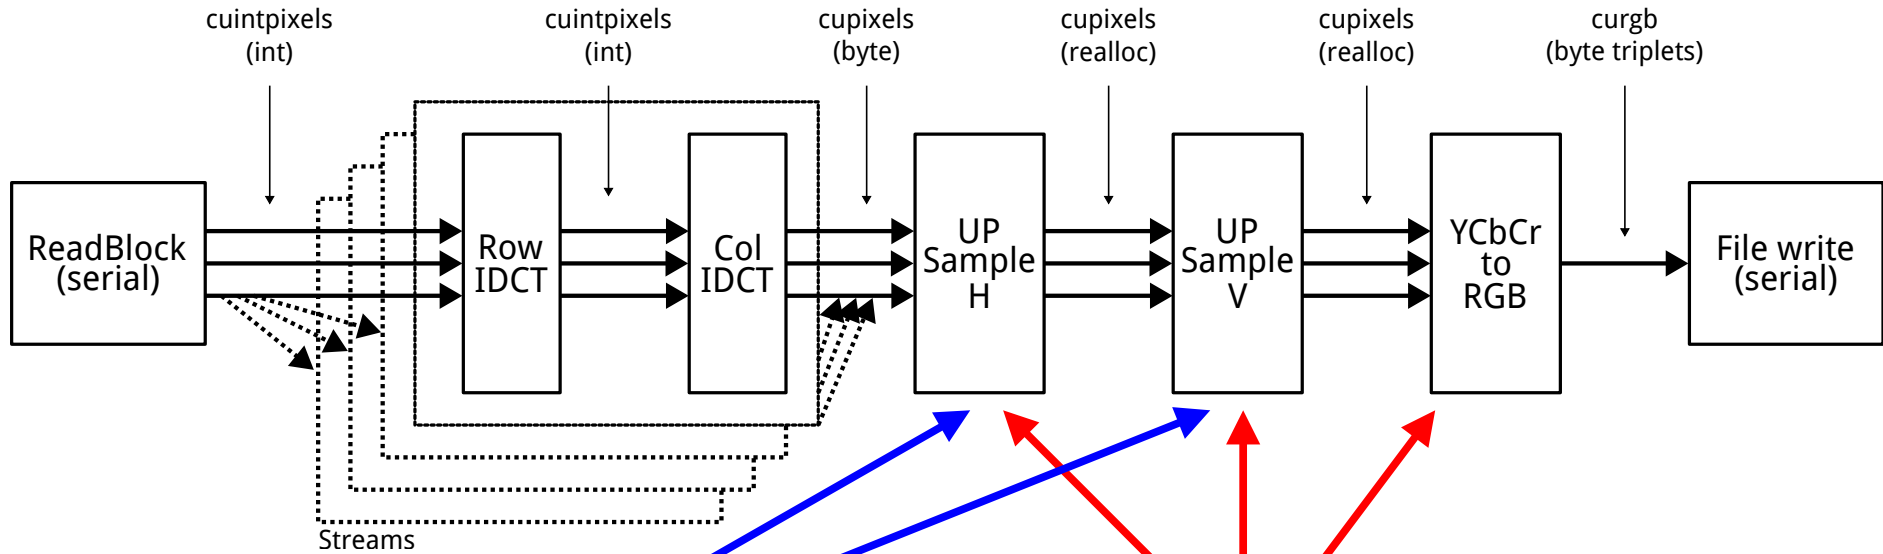


- Few CUDA cores: 1 SM / 128 Threads
- Disk I/O bound workload

Inverse Discrete Cosine Transform (IDCT)



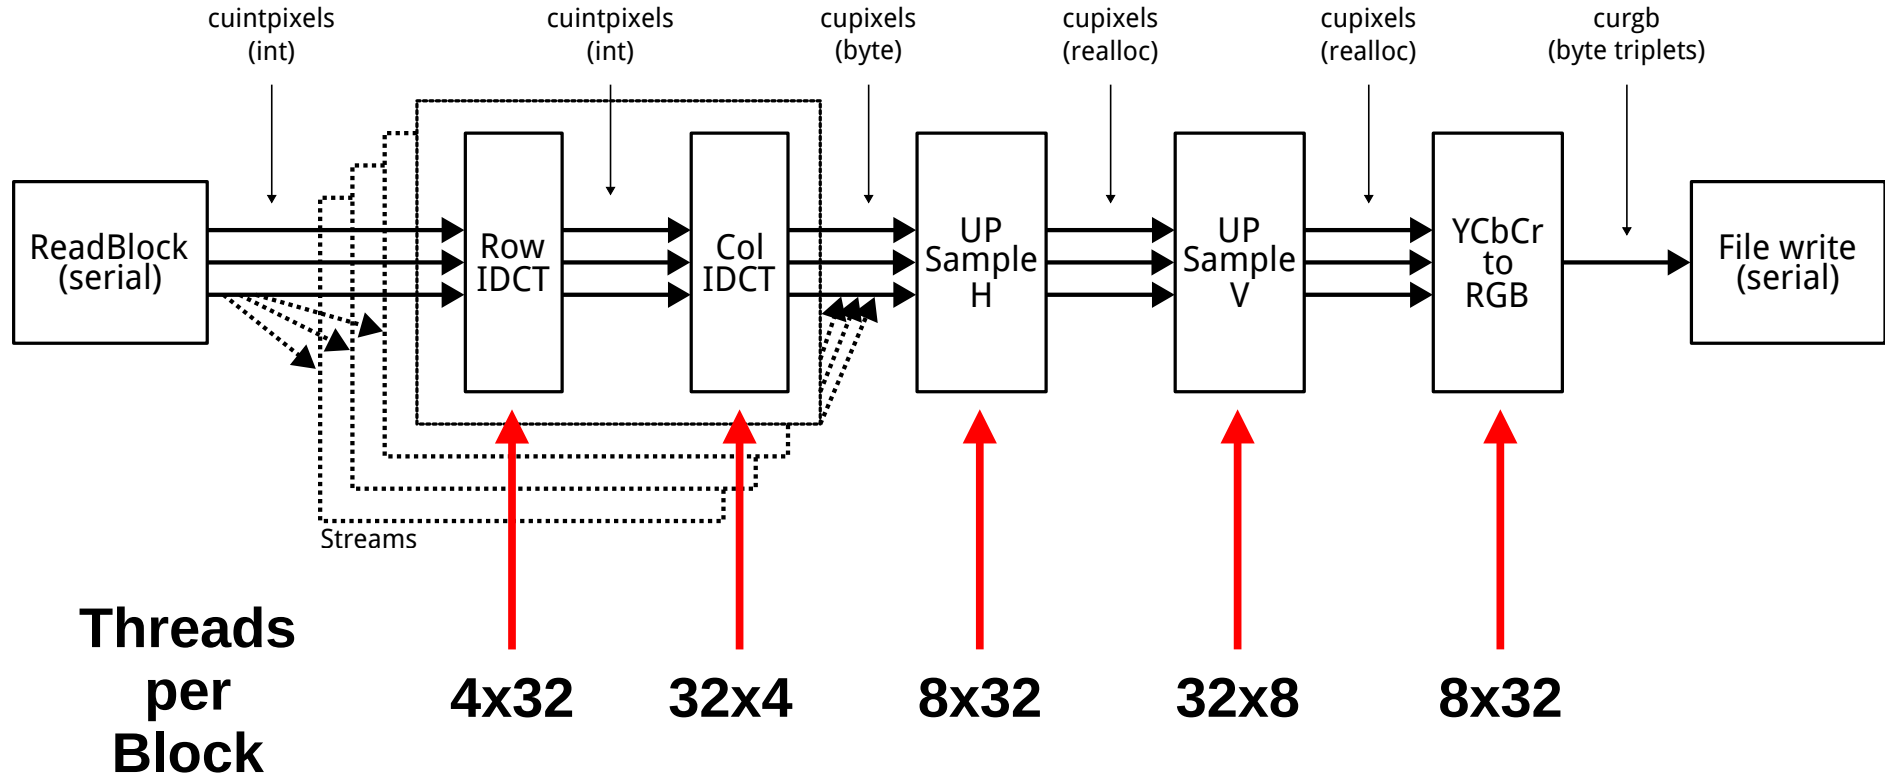
Upsampling, RGB Conversion



**Interdependent
pixels**

**8/16 pixel in per thread
16 pixel out per thread**

CUDA Grids



NanoJPEG CUDA Performance

- Advantageous for big images (on Jetson Nano)
 - 1 Megapixel: GPU **on par** with single-thread CPU
 - 8 Megapixel: GPU **40% faster** than single-thread CPU
- Pro: elaboration in parallel with disk read
- Problems: elaboration is faster than disk I/O (bottleneck)

Performance must be evaluated
for any particular CPU + GPU + DISK combination